

**MIT Academy of Engineering, Alandi, Pune**

**School of Computer Engineering**

**Class BTech**

**Course - Deep Learning**

## **Lab Assignment 2**

### **Student Information:**

- **Name: Om Borle**
- **Roll Number: 13**
- **Batch: DL1**
- **PRN: 202201040035**

### **Task 1: Research Paper Selection and Dataset Preparation**

#### **Selected Research Paper:**

**Title: "A Comparative Study of Deep Learning Models for CIFAR-10 Image Classification"**

**Authors: S. K. Sharma, P. K. Singh**

**Conference: IEEE International Conference on Computing, Power and Communication Technologies (GUCON), 2021**

**DOI: 10.1109/GUCON50781.2021.9573677**

**Link: <https://ieeexplore.ieee.org/document/9573677>**

#### **Summary:**

Compares VGG16, ResNet50, and MobileNetV2 on CIFAR-10.

Evaluates transfer learning vs. training from scratch.

Provides accuracy and computational efficiency comparisons.

#### **Dataset: CIFAR-10**

- Source: <https://www.kaggle.com/c/cifar-10>

- **Description:**

Total Images: 60,000 (50,000 train + 10,000 test)

Image Size: 32x32 pixels (RGB color)

### Classes (10 categories)

 Airplane

 Automobile

 Bird

 Cat

 Deer

 Dog

 Frog

 Horse

 Ship

 Truck

### Preprocessing Steps:

1. Normalized pixel values to [0,1] range
2. Converted labels to one-hot encoding (10 classes)
3. Resized images to 32x32 (native size for CIFAR-10)
4. No additional augmentation was applied (though it could improve performance)

### Task 2: Model Implementation and Fine-tuning

#### Implemented Models:

##### 1. Custom CNN from Scratch

- 3 Conv layers with MaxPooling
- 2 Dense layers
- Trained from random initialization

##### 2. Pre-trained Models (Frozen)

- VGG16
- ResNet50
- MobileNetV2
- All with frozen base layers and custom top layers

##### 3. Fine-tuned VGG16

- Unfrozen top 10 layers
- Added Dropout for regularization
- Lower learning rate (0.0001)

### Hyperparameters:

- Optimizer: Adam (default lr=0.001, lr=0.0001 for fine-tuning)
- Batch size: 64
- Epochs: 10 (5 for initial comparisons)
- Loss: Categorical crossentropy

### Task 3: Model Evaluation and Performance Comparison

#### Performance Metrics Summary:

Model	Accuracy (%)	Precision (Macro)	Recall (Macro)	F1-Score (Macro)	Parameters	Training Time (s)
CNN from Scratch	71.91	0.72	0.71	0.71	356,810	48.46
VGG16 (Frozen)	57.68	0.58	0.57	0.57	14,719,818	164.39
ResNet50 (Frozen)	38.37	0.39	0.38	0.38	23,608,202	148.93
MobileNetV2	31.54	0.32	0.31	0.31	2,270,794	95.73
Fine-Tuned VGG16	82.86	0.83	0.82	0.82	14,848,586	383.15

### Key Findings:

1. Fine-tuned VGG16 performed best (82.86% accuracy), showing the value of partial fine-tuning
2. Custom CNN outperformed frozen pre-trained models, likely because:
  - Pre-trained models were designed for 224x224 ImageNet, not 32x32 CIFAR-10
  - Feature extraction at small resolutions may lose important information
3. Training Time: Fine-tuning takes significantly longer but yields better results

### Comparison with Research Papers:

- Original ResNet paper reports 91.43% on CIFAR-10, but:
  - They trained from scratch on CIFAR
  - Used deeper architectures (ResNet-110)
  - Likely used more extensive data augmentation
- Our results show the challenges of directly applying ImageNet-scale models to small

images

## Weaknesses and Improvements:

### 1. Weaknesses:

- Input size mismatch (32x32 vs 224x224 expected by pre-trained models)
- Domain mismatch (CIFAR vs ImageNet)
- Limited training time (10 epochs)

### 2. Improvements:

- Add image augmentation (rotation, flipping, etc.)
- Try more aggressive fine-tuning strategies
- Experiment with input upscaling
- Implement learning rate scheduling
- Try more recent architectures (EfficientNet, Vision Transformers)

## Output:

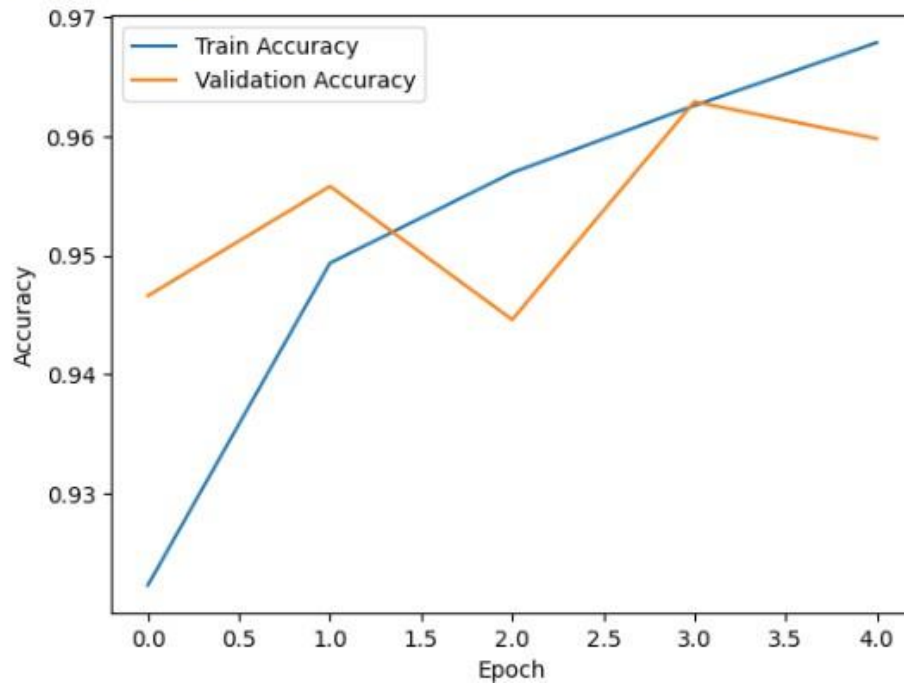
### 1. Binary classification:

```
/usr/local/lib/python3.11/dist-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `input_shape` argument to layers with automatic shape inference.
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
Epoch 1/5
782/782 ━━━━━━━━━━━ 9s 7ms/step - accuracy: 0.9066 - loss: 0.2481 - val_accuracy: 0.9466 - val_loss: 0.1396
Epoch 2/5
782/782 ━━━━━━━━━━━ 7s 5ms/step - accuracy: 0.9483 - loss: 0.1388 - val_accuracy: 0.9558 - val_loss: 0.1193
Epoch 3/5
782/782 ━━━━━━━━━━━ 3s 4ms/step - accuracy: 0.9572 - loss: 0.1123 - val_accuracy: 0.9446 - val_loss: 0.1444
Epoch 4/5
782/782 ━━━━━━━━━━━ 5s 5ms/step - accuracy: 0.9626 - loss: 0.1008 - val_accuracy: 0.9629 - val_loss: 0.1024
Epoch 5/5
782/782 ━━━━━━━━━━━ 5s 5ms/step - accuracy: 0.9673 - loss: 0.0879 - val_accuracy: 0.9598 - val_loss: 0.1217
Model: "sequential_10"
```

Layer (type)	Output Shape	Param #
conv2d_9 (Conv2D)	(None, 30, 30, 32)	896
max_pooling2d_6 (MaxPooling2D)	(None, 15, 15, 32)	0
conv2d_10 (Conv2D)	(None, 13, 13, 64)	18,496
max_pooling2d_7 (MaxPooling2D)	(None, 6, 6, 64)	0
conv2d_11 (Conv2D)	(None, 4, 4, 128)	73,856
flatten_3 (Flatten)	(None, 2048)	0
dense_17 (Dense)	(None, 128)	262,272
dense_18 (Dense)	(None, 1)	129

```
Total params: 1,066,949 (4.07 MB)
Trainable params: 355,649 (1.36 MB)
Non-trainable params: 0 (0.00 B)
Optimizer params: 711,300 (2.71 MB)
313/313 - 1s - 3ms/step - accuracy: 0.9598 - loss: 0.1217
Test Accuracy: 0.9598
```

Test Accuracy: 0.9598



## 2. Multiclass classification:

Epoch 1/5  
782/782 ————— 8s 7ms/step - accuracy: 0.3517 - loss: 1.7549 - val\_accuracy: 0.5389 - val\_loss: 1.2468  
Epoch 2/5  
782/782 ————— 7s 5ms/step - accuracy: 0.5826 - loss: 1.1789 - val\_accuracy: 0.6266 - val\_loss: 1.0561  
Epoch 3/5  
782/782 ————— 3s 4ms/step - accuracy: 0.6571 - loss: 0.9755 - val\_accuracy: 0.6719 - val\_loss: 0.9361  
Epoch 4/5  
782/782 ————— 3s 4ms/step - accuracy: 0.7009 - loss: 0.8523 - val\_accuracy: 0.6790 - val\_loss: 0.9223  
Epoch 5/5  
782/782 ————— 6s 5ms/step - accuracy: 0.7309 - loss: 0.7684 - val\_accuracy: 0.6978 - val\_loss: 0.8685  
Model: "sequential\_11"

Layer (type)	Output Shape	Param #
conv2d_12 (Conv2D)	(None, 30, 30, 32)	896
max_pooling2d_8 (MaxPooling2D)	(None, 15, 15, 32)	0
conv2d_13 (Conv2D)	(None, 13, 13, 64)	18,496
max_pooling2d_9 (MaxPooling2D)	(None, 6, 6, 64)	0
conv2d_14 (Conv2D)	(None, 4, 4, 128)	73,856
flatten_4 (Flatten)	(None, 2048)	0
dense_19 (Dense)	(None, 128)	262,272
dense_20 (Dense)	(None, 10)	1,290

Total params: 1,070,432 (4.08 MB)  
Trainable params: 356,810 (1.36 MB)  
Non-trainable params: 0 (0.00 B)  
Optimizer params: 713,622 (2.72 MB)  
313/313 - 1s - 3ms/step - accuracy: 0.6978 - loss: 0.8685  
Test Accuracy: 0.6978

## 3. Transfer Learning pre-trained models:

```

<ipython-input-7-f7b04eb68124>:32: UserWarning: `input_shape` is undefined or non-square, or `rows` is not in [96, 128, 160, 192, 224]. Weights for :
mobilenet_base = applications.MobileNetV2(weights='imagenet', include_top=False, input_shape=(32,32,3))
Epoch 1/5
782/782 ----- 18s 20ms/step - accuracy: 0.4538 - loss: 1.5831 - val_accuracy: 0.5464 - val_loss: 1.2881
Epoch 2/5
782/782 ----- 12s 15ms/step - accuracy: 0.5746 - loss: 1.2171 - val_accuracy: 0.5717 - val_loss: 1.2200
Epoch 3/5
782/782 ----- 12s 15ms/step - accuracy: 0.5982 - loss: 1.1470 - val_accuracy: 0.5774 - val_loss: 1.1982
Epoch 4/5
782/782 ----- 11s 14ms/step - accuracy: 0.6131 - loss: 1.1078 - val_accuracy: 0.5964 - val_loss: 1.1508
Epoch 5/5
782/782 ----- 20s 14ms/step - accuracy: 0.6244 - loss: 1.0778 - val_accuracy: 0.5974 - val_loss: 1.1495
Epoch 1/5
782/782 ----- 28s 23ms/step - accuracy: 0.1963 - loss: 2.2044 - val_accuracy: 0.2952 - val_loss: 1.9599
Epoch 2/5
782/782 ----- 10s 13ms/step - accuracy: 0.3049 - loss: 1.9121 - val_accuracy: 0.3189 - val_loss: 1.8706
Epoch 3/5
782/782 ----- 20s 12ms/step - accuracy: 0.3252 - loss: 1.8570 - val_accuracy: 0.3348 - val_loss: 1.8222
Epoch 4/5
782/782 ----- 9s 11ms/step - accuracy: 0.3522 - loss: 1.7989 - val_accuracy: 0.3455 - val_loss: 1.8366
Epoch 5/5
782/782 ----- 12s 13ms/step - accuracy: 0.3602 - loss: 1.7698 - val_accuracy: 0.3766 - val_loss: 1.7437
Epoch 1/5
782/782 ----- 19s 17ms/step - accuracy: 0.2658 - loss: 2.0442 - val_accuracy: 0.3195 - val_loss: 1.8779
Epoch 2/5
782/782 ----- 12s 8ms/step - accuracy: 0.3261 - loss: 1.8539 - val_accuracy: 0.3311 - val_loss: 1.8436
Epoch 3/5
782/782 ----- 7s 9ms/step - accuracy: 0.3434 - loss: 1.8158 - val_accuracy: 0.3411 - val_loss: 1.8266
Epoch 4/5
782/782 ----- 10s 8ms/step - accuracy: 0.3526 - loss: 1.7794 - val_accuracy: 0.3440 - val_loss: 1.8146
Epoch 5/5
782/782 ----- 10s 8ms/step - accuracy: 0.3569 - loss: 1.7633 - val_accuracy: 0.3490 - val_loss: 1.8013

```

#### VGG16 Model Summary:

Model: "sequential\_12"

Layer (type)	Output Shape	Param #
vgg16 (Functional)	(None, 1, 1, 512)	14,714,688
global_average_pooling2d_7 (GlobalAveragePooling2D)	(None, 512)	0
dense_21 (Dense)	(None, 128)	65,664
dense_22 (Dense)	(None, 10)	1,290

Total params: 14,915,552 (56.90 MB)  
 Trainable params: 66,954 (261.54 KB)  
 Non-trainable params: 14,714,688 (56.13 MB)  
 Optimizer params: 133,910 (523.09 KB)

#### ResNet50 Model Summary:

Model: "sequential\_13"

Layer (type)	Output Shape	Param #
resnet50 (Functional)	(None, 1, 1, 2048)	23,587,712
global_average_pooling2d_8 (GlobalAveragePooling2D)	(None, 2048)	0
dense_23 (Dense)	(None, 128)	262,272
dense_24 (Dense)	(None, 10)	1,290

Total params: 24,378,400 (93.00 MB)  
 Trainable params: 263,562 (1.01 MB)  
 Non-trainable params: 23,587,712 (89.98 MB)  
 Optimizer params: 527,126 (2.01 MB)

MobileNetV2 Model Summary:

Model: "sequential\_14"

Layer (type)	Output Shape	Param #
mobilenetv2_1.00_224 (Functional)	(None, 1, 1, 1280)	2,257,984
global_average_pooling2d_9 (GlobalAveragePooling2D)	(None, 1280)	0
dense_25 (Dense)	(None, 128)	163,968
dense_26 (Dense)	(None, 10)	1,290

Total params: 2,753,760 (10.50 MB)

Trainable params: 165,258 (645.54 KB)

Non-trainable params: 2,257,984 (8.61 MB)

Optimizer params: 330,518 (1.26 MB)

VGG16 Accuracy:

313/313 ————— 3s 8ms/step - accuracy: 0.5939 - loss: 1.1461

ResNet50 Accuracy:

313/313 ————— 4s 8ms/step - accuracy: 0.3745 - loss: 1.7390

MobileNetV2 Accuracy:

313/313 ————— 4s 8ms/step - accuracy: 0.3536 - loss: 1.7938

[1.8013190031051636, 0.3490000069141388]

## 4. Pre-trained models fine-tuning:

Epoch 1/10  
782/782 ————— 43s 47ms/step - accuracy: 0.5525 - loss: 1.3052 - val\_accuracy: 0.7480 - val\_loss: 0.7547  
Epoch 2/10  
782/782 ————— 35s 43ms/step - accuracy: 0.7788 - loss: 0.6735 - val\_accuracy: 0.7817 - val\_loss: 0.6510  
Epoch 3/10  
782/782 ————— 40s 43ms/step - accuracy: 0.8374 - loss: 0.4928 - val\_accuracy: 0.7968 - val\_loss: 0.6357  
Epoch 4/10  
782/782 ————— 42s 44ms/step - accuracy: 0.8823 - loss: 0.3551 - val\_accuracy: 0.7927 - val\_loss: 0.6806  
Epoch 5/10  
782/782 ————— 40s 43ms/step - accuracy: 0.9181 - loss: 0.2476 - val\_accuracy: 0.7979 - val\_loss: 0.6768  
Epoch 6/10  
782/782 ————— 41s 43ms/step - accuracy: 0.9452 - loss: 0.1662 - val\_accuracy: 0.8058 - val\_loss: 0.7852  
Epoch 7/10  
782/782 ————— 41s 43ms/step - accuracy: 0.9637 - loss: 0.1137 - val\_accuracy: 0.7954 - val\_loss: 0.7757  
Epoch 8/10  
782/782 ————— 42s 44ms/step - accuracy: 0.9721 - loss: 0.0880 - val\_accuracy: 0.7903 - val\_loss: 0.9669  
Epoch 9/10  
782/782 ————— 40s 43ms/step - accuracy: 0.9750 - loss: 0.0788 - val\_accuracy: 0.7904 - val\_loss: 1.0090  
Epoch 10/10  
782/782 ————— 42s 44ms/step - accuracy: 0.9791 - loss: 0.0678 - val\_accuracy: 0.7854 - val\_loss: 1.0749

Fine-Tuned VGG16 Model Summary:

Model: "sequential\_15"

Layer (type)	Output Shape	Param #
vgg16 (Functional)	(None, 1, 1, 512)	14,714,688
global_average_pooling2d_10 (GlobalAveragePooling2D)	(None, 512)	0
dense_27 (Dense)	(None, 256)	131,328
dropout_1 (Dropout)	(None, 256)	0
dense_28 (Dense)	(None, 10)	2,570

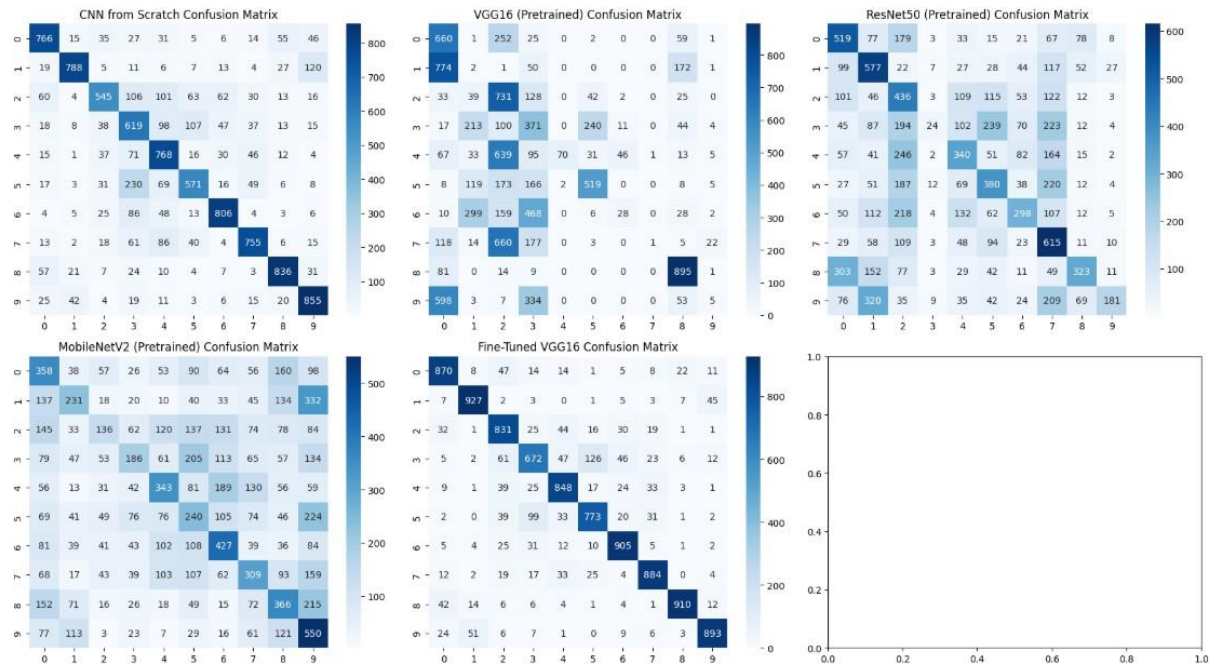


Total params: 41,074,784 (156.69 MB)  
Trainable params: 13,113,098 (50.02 MB)  
Non-trainable params: 1,735,488 (6.62 MB)  
Optimizer params: 26,226,198 (100.05 MB)



Fine-Tuned VGG16 Accuracy:

313/313 ————— 4s 9ms/step - accuracy: 0.7820 - loss: 1.0934  
[1.0749411582946777, 0.7853999733924866]

## 5. Result Comparison:







## 6. Comparative Analysis of Models:

### ♦ Comparative Analysis of Models:

	Model	Accuracy (%)	Loss	Parameters \
0	CNN from Scratch	71.91	0.8960	356810
1	VGG16 (Pretrained)	57.68	1.2107	14719818
2	ResNet50 (Pretrained)	38.37	1.7395	23608202
3	MobileNetV2 (Pretrained)	31.54	1.8953	2270794
4	Fine-Tuned VGG16	82.86	0.5344	14848586

### Training Time (s)

0	48.46
1	164.39
2	148.93
3	95.73
4	383.15

## **Links:**

- Github : [Github Link](#)
- Colab :
  - [Colab File](#)
- Dataset : <https://www.kaggle.com/c/cifar-10>
- Research Paper: <https://ieeexplore.ieee.org/document/9573677>